

MAWLANA BHASHANI SCIENCE AND TECHNOLOGY UNIVERSITY

Santosh,Tangail – 1902



Course Title : Computer Networks Lab

**Lab Report : SDN Controllers and Mininet
Name**

Lab Report No. : 07

Submitted by,

Name : ANIKA JAHIN

ID: IT-17056

Session: 2016-17

Dept. of ICT,MBSTU.

Submitted to,

NAZRUL ISLAM

Assistant Professor

Dept. of ICT,MBSTU.

Theory:

Traffic Generator:

What is iPerf?: iPerf is a tool for active measurements of the maximum achievable bandwidth on IP networks. It supports tuning of various parameters related to timing, buffers and protocols (TCP, UDP, SCTP with IPv4 and IPv6). For each test it reports the bandwidth, loss, and other parameters.

Mininet: Mininet creates a **realistic virtual network**, running **real kernel, switch and application code**, on a single machine (VM, cloud or native) Because you can easily interact with your network using the Mininet CLI (and API), customize it, share it with others, or deploy it on real hardware, Mininet is useful for development, teaching, and research. Mininet is also a great way to develop, share, and experiment with OpenFlow and Software-Defined Networking systems.

Install iperf:

```
anika@anika-VirtualBox:~$ sudo apt-get install iperf
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  iperf
0 upgraded, 1 newly installed, 0 to remove and 388 not upgraded.
Need to get 60.5 kB of archives.
After this operation, 176 kB of additional disk space will be used.
Get:1 http://bd.archive.ubuntu.com/ubuntu bionic-updates/universe amd64 iperf a
md64 2.0.10+dfsg1-1ubuntu0.18.04.2 [60.5 kB]
Fetched 60.5 kB in 2s (34.5 kB/s)
Selecting previously unselected package iperf.
(Reading database ... 131970 files and directories currently installed.)
Preparing to unpack .../iperf_2.0.10+dfsg1-1ubuntu0.18.04.2_amd64.deb ...
Unpacking iperf (2.0.10+dfsg1-1ubuntu0.18.04.2) ...
Setting up iperf (2.0.10+dfsg1-1ubuntu0.18.04.2) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
anika@anika-VirtualBox:~$
```

Install Mininet:

```

anika@anika-VirtualBox:~$ sudo apt-get install mininet
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  cgroup-bin cgroup-tools libcgroup1 openvswitch-common openvswitch-switch
  python-pkg-resources python-six
Suggested packages:
  openvswitch-doc python-setuptools
The following NEW packages will be installed:
  cgroup-bin cgroup-tools libcgroup1 mininet openvswitch-common
  openvswitch-switch python-pkg-resources python-six
0 upgraded, 8 newly installed, 0 to remove and 388 not upgraded.

```

4. Exercises

Exercise 4.1.1: Open a Linux terminal, and execute the command line *iperf --help*. Provide four configuration options of iperf.

```

anika@anika-VirtualBox:~$ iperf --help
Usage: iperf [-s|-c host] [options]
       iperf [-h|--help] [-v|--version]

Client/Server:
  -b, --bandwidth #[kmgKMG | pps]  bandwidth to send at in bits/sec or packets
per second
  -e, --enhancedreports             use enhanced reporting giving more tcp/udp and traff
ic information
  -f, --format [kmgKMG]            format to report: Kbits, Mbits, KBytes, MBytes
  -i, --interval #                 seconds between periodic bandwidth reports
  -l, --len #[kmKM]                length of buffer in bytes to read or write (Defaul
ts: TCP=128K, v4 UDP=1470, v6 UDP=1450)
  -m, --print_mss                  print TCP maximum segment size (MTU - TCP/IP header)
  -o, --output <filename>          output the report or error message to this specifi
ed file
  -p, --port #                     server port to listen on/connect to
  -u, --udp                        use UDP rather than TCP
      --udp-counters-64bit          use 64 bit sequence numbers with UDP
  -w, --window #[KM]              TCP window size (socket buffer size)
  -z, --realtime                   request realtime scheduler
  -B, --bind <host>               bind to <host>, an interface or multicast address
  -C, --compatibility              for use with older versions does not sent extra msgs
  -M, --mss #                     set TCP maximum segment size (MTU - 40 bytes)
  -N, --nodelay                   set TCP no delay, disabling Nagle's Algorithm
  -S, --tos #                     set the socket's IP_TOS (byte) field

```

Exercise 4.1.2: Open two Linux terminals, and configure terminal-1 as client (*iperf -c IPv4_server_address*) and terminal-2 as server (*iperf -s*).

For terminal -2:

```
anika@anika-VirtualBox:~$ iperf -c 127.0.0.1
-----
Client connecting to 127.0.0.1, TCP port 5001
TCP window size: 2.50 MByte (default)
-----
[  3] local 127.0.0.1 port 54292 connected with 127.0.0.1 port 5001
[ ID] Interval      Transfer    Bandwidth
[  3]  0.0-10.0 sec  27.8 GBytes 23.9 Gbits/sec
anika@anika-VirtualBox:~$
```

For terminal -1:

```
anika@anika-VirtualBox:~$ iperf -s
-----
Server listening on TCP port 5001
TCP window size: 128 KByte (default)
-----
[  4] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 54292
[ ID] Interval      Transfer    Bandwidth
[  4]  0.0-10.0 sec  27.8 GBytes 23.9 Gbits/sec
```

Exercise 4.1.3: Open two Linux terminals, and configure terminal-1 as client and terminal-2 as server for exchanging UDP traffic, which are the command lines? Which are the statistics are provided at the end of transmission?

```
anika@anika-VirtualBox:~$ iperf -c 127.0.0.1 -u
-----
Client connecting to 127.0.0.1, UDP port 5001
Sending 1470 byte datagrams, IPG target: 11215.21 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[  3] local 127.0.0.1 port 46294 connected with 127.0.0.1 port 5001
[ ID] Interval      Transfer    Bandwidth
[  3]  0.0-10.0 sec  1.25 MBytes 1.05 Mbits/sec
[  3] Sent 893 datagrams
[  3] Server Report:
[  3]  0.0-10.0 sec  1.25 MBytes 1.05 Mbits/sec  0.000 ms  0/ 893 (0%)
anika@anika-VirtualBox:~$
```

```

anika@anika-VirtualBox:~$ iperf -s -u
-----
Server listening on UDP port 5001
Receiving 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 46294
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagram
s
[ 3] 0.0-10.0 sec  1.25 MBytes  1.05 Mbits/sec  0.002 ms   0/ 893 (0%)

[ 3] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 46294
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagram
s
[ 3] 0.0-10.0 sec  1.25 MBytes  1.05 Mbits/sec  0.002 ms   0/ 893 (0%)

```

Exercise 4.1.4: Open two Linux terminals, and configure terminal-1 as client and terminal-2 as server for exchanging UDP traffic, with:

- o Packet length = 1000bytes
- o Time = 20 seconds
- o Bandwidth = 1Mbps
- o Port = 9900

Which are the command lines?

The command lines are:

For terminal 1:

`iperf -c 127.0.0.1 -u -l 1000 -t 20 -b 1 -p 9900`

```

anika@anika-VirtualBox:~$ iperf -c 127.0.0.1 -u -l 1000 -t 20 -b 1 -p 9900
WARNING: delay too large, reducing from 8000.0 to 1.0 seconds.
-----
Client connecting to 127.0.0.1, UDP port 9900
Sending 1000 byte datagrams, IPG target: 8000000000.00 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 3] local 127.0.0.1 port 37049 connected with 127.0.0.1 port 9900
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-21.0 sec  20.5 KBytes  8.00 Kbits/sec
[ 3] Sent 21 datagrams
[ 3] Server Report:
[ 3] 0.0-21.0 sec  20.5 KBytes  8.00 Kbits/sec  0.000 ms   0/ 21 (0%)
anika@anika-VirtualBox:~$

```

For terminal 2:

`lperf -s -u -p 9900`

```
anika@anika-VirtualBox:~$ iperf -s -u -p 9900
-----
Server listening on UDP port 9900
Receiving 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 127.0.0.1 port 9900 connected with 127.0.0.1 port 37049
[ ID] Interval      Transfer      Bandwidth      Jitter    Lost/Total Datagram
s
[ 3]  0.0-21.0 sec  20.5 KBytes   8.00 Kbits/sec  62.509 ms   0/ 21 (0%)
```

Using Mininet

Exercise 4.2.1: Open two Linux terminals, and execute the command line *ifconfig* in terminal-1. How many interfaces are present?

In terminal-2, execute the command line *sudo mn*, which is the output?

In terminal-1 execute the command line *ifconfig*. How many real and virtual interfaces are present now?

```
anika@anika-VirtualBox:~$ ifconfig
enp0s3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.2.15 netmask 255.255.255.0 broadcast 10.0.2.255
    inet6 fe80::d04d:cb3e:2101:b7ed prefixlen 64 scopeid 0x20<link>
    ether 08:00:27:aa:2c:bc txqueuelen 1000 (Ethernet)
    RX packets 3922 bytes 3713725 (3.7 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 2459 bytes 199091 (199.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 718920 bytes 29904911145 (29.9 GB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 718920 bytes 29904911145 (29.9 GB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

anika@anika-VirtualBox:~$
```

```
anika@anika-VirtualBox:~$ sudo mn
[sudo] password for anika:
*** No default OpenFlow controller found for default switch!
*** Falling back to OVS Bridge
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1)
*** Configuring hosts
h1 h2
*** Starting controller

*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet>
```

```
anika@anika-VirtualBox:~$ ifconfig
enp0s3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.2.15 netmask 255.255.255.0 broadcast 10.0.2.255
    inet6 fe80::d04d:cb3e:2101:b7ed prefixlen 64 scopeid 0x20<link>
    ether 08:00:27:aa:2c:bc txqueuelen 1000 (Ethernet)
    RX packets 3924 bytes 3713904 (3.7 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 2461 bytes 199238 (199.2 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 718932 bytes 29904912117 (29.9 GB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 718932 bytes 29904912117 (29.9 GB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::70c9:2cff:fef7:f8ce prefixlen 64 scopeid 0x20<link>
    ether 72:c9:2c:f7:f8:ce txqueuelen 1000 (Ethernet)
    RX packets 10 bytes 796 (796.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 32 bytes 3571 (3.5 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```



```
s1-eth2: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::5c10:9ff:fed4:a8e6 prefixlen 64 scopeid 0x20<link>
    ether 5e:10:09:d4:a8:e6 txqueuelen 1000 (Ethernet)
RX packets 10 bytes 796 (796.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 33 bytes 3661 (3.6 KB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

Exercise 4.2.2: Interacting with mininet; in terminal-2, display the following command lines and explain what it does:

mininet> help

```
mininet> help

Documented commands (type help <topic>):
=====
EOF      gterm   iperfudp nodes      pingpair  py        switch
dpctl    help    link     noecho     pingpairfull quit      time
dump     intfs   links    pingall    ports     sh        x
exit     iperf   net      pingallfull px         source    xterm

You may also send a command to a node using:
    <node> command {args}
For example:
    mininet> h1 ifconfig

The interpreter automatically substitutes IP addresses
for node names when a node is the first arg, so commands
like
    mininet> h2 ping h3
should work.

Some character-oriented interactive commands require
noecho:
    mininet> noecho h2 vi foo.py
However, starting up an xterm/gterm is generally better:
    mininet> xterm h2
```

mininet> nodes

```
mininet> nodes
available nodes are:
h1 h2 s1
mininet> █
```


mininet> net

```
mininet> net
h1 h1-eth0:s1-eth1
h2 h2-eth0:s1-eth2
s1 lo: s1-eth1:h1-eth0 s1-eth2:h2-eth0
mininet>
```

mininet> dump

```
mininet> dump
<Host h1: h1-eth0:10.0.0.1 pid=4532>
<Host h2: h2-eth0:10.0.0.2 pid=4534>
<OVSBridge s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None pid=4539>
mininet>
```

mininet> h1 ifconfig -a

```
mininet> h1 ifconfig -a
h1-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.1 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::4024:79ff:fe36:8015 prefixlen 64 scopeid 0x20<link>
    ether 42:24:79:36:80:15 txqueuelen 1000 (Ethernet)
    RX packets 40 bytes 4242 (4.2 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 13 bytes 1006 (1.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

mininet>
```

mininet> s1 ifconfig -a

```
mininet> s1 ifconfig -a
enp0s3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.2.15 netmask 255.255.255.0 broadcast 10.0.2.255
    inet6 fe80::d04d:cb3e:2101:b7ed prefixlen 64 scopeid 0x20<link>
    ether 08:00:27:aa:2c:bc txqueuelen 1000 (Ethernet)
    RX packets 3952 bytes 3719718 (3.7 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 2486 bytes 201960 (201.9 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 718944 bytes 29904913245 (29.9 GB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 718944 bytes 29904913245 (29.9 GB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ovs-system: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether d2:3a:6c:20:10:e6 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether b2:6b:4e:3e:76:4c txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
```

```

s1: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether b2:6b:4e:3e:76:4c txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 21 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::70c9:2cff:fe7:f8ce prefixlen 64 scopeid 0x20<link>
    ether 72:c9:2c:f7:f8:ce txqueuelen 1000 (Ethernet)
    RX packets 13 bytes 1006 (1.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 41 bytes 4312 (4.3 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1-eth2: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::5c10:9ff:fed4:a8e6 prefixlen 64 scopeid 0x20<link>
    ether 5e:10:09:d4:a8:e6 txqueuelen 1000 (Ethernet)
    RX packets 13 bytes 1006 (1.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 42 bytes 4402 (4.4 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

mininet> h1 ping -c 5 h2

```

mininet> h1 ping -c 5 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.388 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.115 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.145 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.132 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.130 ms

--- 10.0.0.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4097ms
rtt min/avg/max/mdev = 0.115/0.182/0.388/0.103 ms
mininet>

```

Exercise 4.2.3: In terminal-2, display the following command line: *sudo mn --link tc,bw=10,delay=500ms*

o mininet> h1 ping -c 5 h2, What happen with the link?

o mininet> h1 iperf -s -u &

o mininet> h2 iperf -c IPv4_h1 -u, Is there any packet loss?

```

anika@anika-VirtualBox:~$ sudo mn --link tc,bw=10,delay=500ms
[sudo] password for anika:
*** No default OpenFlow controller found for default switch!
**Thunderbird Mailk to OVS Bridge
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
-----
-
Caught exception. Cleaning up...

Exception: Error creating interface pair (h1-eth0,s1-eth1): RTNETLINK answers:
File exists

-----
-
*** Removing excess controllers/ofprotocols/ofdatapaths/pings/noxes
killall controller ofprotocol ofdatapath ping nox_core lt-nox_core ovs-openflow
d ovs-controller udpbwtest mnexec ivs 2> /dev/null
killall -9 controller ofprotocol ofdatapath ping nox_core lt-nox_core ovs-openf
lowd ovs-controller udpbwtest mnexec ivs 2> /dev/null
pkill -9 -f "sudo mnexec"
*** Removing junk from /tmp
rm -f /tmp/vconn* /tmp/vlogs* /tmp/*.out /tmp/*.log
*** Removing old X11 tunnels

```

```

mininet> h1 ping -c 5 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.388 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.115 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.145 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.132 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.130 ms

--- 10.0.0.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4097ms
rtt min/avg/max/mdev = 0.115/0.182/0.388/0.103 ms
mininet> h1 ping -c 5 h2

```

Conclusion: Mininet is a useful tool for teaching, development and research. With it, a realistic virtual network, running a real kernel switch and application code, can be set up in a few seconds on a single machine, either virtual or native. It is actively developed and supported.

Emulation refers to the running of unchanged code on virtual hardware on the top of the physical host, interactively. It is handy, practical and low cost. It comes with certain restrictions, though, like slower speeds compared to running the same code on a hardware test-bed which

is fast and accurate, but expensive. While a simulator requires code modifications and is slow as well. Mininet is a network emulator that enables the creation of a network of virtual hosts, switches, controllers, and links. Mininet hosts standard Linux network software, and its switches support OpenFlow, a software defined network (SDN) for highly flexible custom routing.

It constructs a virtual network that appears to be a real physical network. You can create a network topology, simulate it and implement the various network performance parameters such as bandwidth, latency, packet loss, etc, with Mininet, using simple code. You can create the virtual network on a single machine (a VM, the cloud or a native machine).

Mininet permits the creation of multiple nodes (hosts, switches or controllers), enabling a big network to be simulated on a single PC. This is very useful in experimenting with various topologies and different controllers, for different network scenarios.

The programs that you run can send packets through virtual switches that seem like real Ethernet interfaces, with a given link speed and delay. Packets get processed by what looks like a real Ethernet switch, router, or middle-box, with a given amount of queuing.

The Mininet CLI and API facilitate easy interaction with our network. Virtual hosts, switches, links and controllers created through Mininet are the real thing. They are just created using the Mininet emulator rather than hardware and for the most part, their behaviour is similar to discrete hardware elements.